

# TikTok Regression modeling

January 8, 2025

## 1 TikTok Project

### Course 5 - Regression Analysis: Simplify complex data relationships

You are a data professional at TikTok. The data team is working towards building a machine learning model that can be used to determine whether a video contains a claim or whether it offers an opinion. With a successful prediction model, TikTok can reduce the backlog of user reports and prioritize them more efficiently.

The team is getting closer to completing the project, having completed an initial plan of action, initial Python coding work, EDA, and hypothesis testing.

The TikTok team has reviewed the results of the hypothesis testing. TikTok's Operations Lead, Maika Abadi, is interested in how different variables are associated with whether a user is verified. Earlier, the data team observed that if a user is verified, they are much more likely to post opinions. Now, the data team has decided to explore how to predict verified status to help them understand how video characteristics relate to verified users. Therefore, you have been asked to conduct a logistic regression using verified status as the outcome variable. The results may be used to inform the final model related to predicting whether a video is a claim vs an opinion.

A notebook was structured and prepared to help you in this project. Please complete the following questions.

## 2 Course 5 End-of-course project: Regression modeling

In this activity, you will build a logistic regression model in Python. As you have learned, logistic regression helps you estimate the probability of an outcome. For data science professionals, this is a useful skill because it allows you to consider more than one variable against the variable you're measuring against. This opens the door for much more thorough and flexible analysis to be completed.

**The purpose** of this project is to demonstrate knowledge of EDA and regression models.

**The goal** is to build a logistic regression model and evaluate the model. *This activity has three parts:*

**Part 1:** EDA & Checking Model Assumptions \* What are some purposes of EDA before constructing a logistic regression model?

**Part 2:** Model Building and Evaluation \* What resources do you find yourself using as you complete this stage?

### Part 3: Interpreting Model Results

- What key insights emerged from your model(s)?
- What business recommendations do you propose based on the models built?

Follow the instructions and answer the question below to complete the activity. Then, you will complete an executive summary using the questions listed on the PACE Strategy Document.

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

## 3 Build a regression model

### 4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

#### 4.1 PACE: Plan

Consider the questions in your PACE Strategy Document to reflect on the Plan stage.

##### 4.1.1 Task 1. Imports and loading

Import the data and packages that you've learned are needed for building regression models.

```
[1]: # Import packages for data manipulation
import pandas as pd
import numpy as np

# Import packages for data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Import packages for data preprocessing
from sklearn.preprocessing import OneHotEncoder
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.utils import resample

# Import packages for data modeling
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```

Load the TikTok dataset.

**Note:** As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed

with this lab. Please continue with this activity by completing the following instructions.

```
[2]: # Load dataset into dataframe
data = pd.read_csv("tiktok_dataset.csv")
```

## 4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document to reflect on the Analyze stage.

In this stage, consider the following question where applicable to complete your code response:

- What are some purposes of EDA before constructing a logistic regression model?

**Answer:** The purposes of EDA before logistic regression are to detect data anomalies (like outliers and class imbalances) and to confirm no severe multicollinearity.

### 4.2.1 Task 2a. Explore data with EDA

Analyze the data and check for and handle missing values and duplicates.

Inspect the first five rows of the dataframe.

```
[3]: # Display first few rows
#### YOUR CODE HERE ####
data.head(5)
```

```
[3]: # claim_status    video_id  video_duration_sec  \
0  1          claim  7017666017             59
1  2          claim  4014381136             32
2  3          claim  9859838091             31
3  4          claim  1866847991             25
4  5          claim  7105231098             19

              video_transcription_text  verified_status  \
0  someone shared with me that drone deliveries a...  not verified
1  someone shared with me that there are more mic...  not verified
2  someone shared with me that american industria...  not verified
3  someone shared with me that the metro of st. p...  not verified
4  someone shared with me that the number of busi...  not verified

author_ban_status  video_view_count  video_like_count  video_share_count  \
0      under review      343296.0      19425.0           241.0
1           active      140877.0      77355.0          19034.0
2           active      902185.0      97690.0           2858.0
3           active      437506.0     239954.0          34812.0
4           active      56167.0      34987.0           4110.0

video_download_count  video_comment_count
0                1.0                0.0
1             1161.0             684.0
```

2	833.0	329.0
3	1234.0	584.0
4	547.0	152.0

Get the number of rows and columns in the dataset.

```
[4]: # Get number of rows and columns
    ### YOUR CODE HERE ###
    data.shape
```

```
[4]: (19382, 12)
```

Get the data types of the columns.

```
[5]: # Get data types of columns
    ### YOUR CODE HERE ###
    data.dtypes
```

```
[5]: #                                     int64
    claim_status                        object
    video_id                            int64
    video_duration_sec                  int64
    video_transcription_text            object
    verified_status                     object
    author_ban_status                   object
    video_view_count                    float64
    video_like_count                    float64
    video_share_count                   float64
    video_download_count                float64
    video_comment_count                 float64
    dtype: object
```

Get basic information about the dataset.

```
[6]: # Get basic information
    ### YOUR CODE HERE ###
    data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19382 entries, 0 to 19381
Data columns (total 12 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   #                                     19382 non-null  int64
 1   claim_status                         19084 non-null  object
 2   video_id                             19382 non-null  int64
 3   video_duration_sec                   19382 non-null  int64
 4   video_transcription_text             19084 non-null  object
 5   verified_status                      19382 non-null  object
```

```

6  author_ban_status      19382 non-null  object
7  video_view_count      19084 non-null  float64
8  video_like_count      19084 non-null  float64
9  video_share_count      19084 non-null  float64
10 video_download_count   19084 non-null  float64
11 video_comment_count    19084 non-null  float64
dtypes: float64(5), int64(3), object(4)
memory usage: 1.8+ MB

```

Generate basic descriptive statistics about the dataset.

```

[7]: # Generate basic descriptive stats
    ### YOUR CODE HERE ###
    data.describe()

```

```

[7]:
count      #      video_id  video_duration_sec  video_view_count  \
count  19382.000000  1.938200e+04      19382.000000      19084.000000
mean    9691.500000  5.627454e+09      32.421732      254708.558688
std     5595.245794  2.536440e+09      16.229967      322893.280814
min       1.000000  1.234959e+09       5.000000       20.000000
25%     4846.250000  3.430417e+09      18.000000      4942.500000
50%     9691.500000  5.618664e+09      32.000000      9954.500000
75%    14536.750000  7.843960e+09      47.000000     504327.000000
max    19382.000000  9.999873e+09      60.000000     999817.000000

      video_like_count  video_share_count  video_download_count  \
count      19084.000000      19084.000000      19084.000000
mean       84304.636030      16735.248323      1049.429627
std       133420.546814      32036.174350      2004.299894
min           0.000000           0.000000           0.000000
25%         810.750000          115.000000           7.000000
50%        3403.500000          717.000000          46.000000
75%       125020.000000      18222.000000       1156.250000
max       657830.000000     256130.000000     14994.000000

      video_comment_count
count      19084.000000
mean         349.312146
std         799.638865
min           0.000000
25%           1.000000
50%           9.000000
75%         292.000000
max        9599.000000

```

Check for and handle missing values.

```
[8]: # Check for missing values
    ### YOUR CODE HERE ###
    data.isna().sum()
```

```
[8]: #
    claim_status      298
    video_id          0
    video_duration_sec 0
    video_transcription_text 298
    verified_status    0
    author_ban_status  0
    video_view_count   298
    video_like_count   298
    video_share_count  298
    video_download_count 298
    video_comment_count 298
    dtype: int64
```

```
[9]: # Drop rows with missing values
    ### YOUR CODE HERE ###
    data = data.dropna(axis = 0)
```

```
[10]: # Display first few rows after handling missing values
    ### YOUR CODE HERE ###
    data.head(5)
```

```
[10]: # claim_status    video_id    video_duration_sec \
0  1      claim    7017666017      59
1  2      claim    4014381136      32
2  3      claim    9859838091      31
3  4      claim    1866847991      25
4  5      claim    7105231098      19

      video_transcription_text    verified_status \
0  someone shared with me that drone deliveries a...    not verified
1  someone shared with me that there are more mic...    not verified
2  someone shared with me that american industria...    not verified
3  someone shared with me that the metro of st. p...    not verified
4  someone shared with me that the number of busi...    not verified

      author_ban_status    video_view_count    video_like_count    video_share_count \
0      under review      343296.0      19425.0      241.0
1      active      140877.0      77355.0      19034.0
2      active      902185.0      97690.0      2858.0
3      active      437506.0      239954.0      34812.0
4      active      56167.0      34987.0      4110.0
```

	video_download_count	video_comment_count
0	1.0	0.0
1	1161.0	684.0
2	833.0	329.0
3	1234.0	584.0
4	547.0	152.0

Check for and handle duplicates.

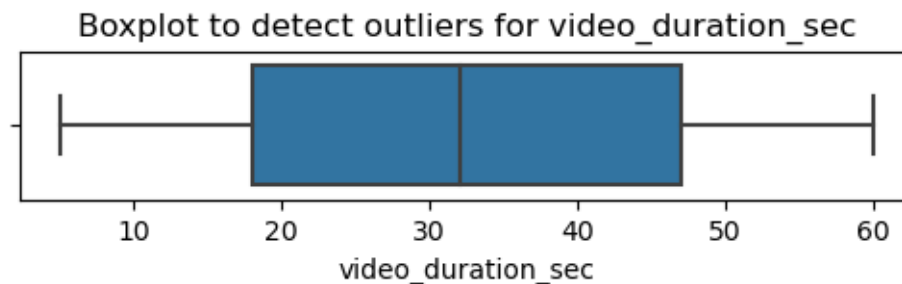
```
[11]: # Check for duplicates
      ### YOUR CODE HERE ###
      data.duplicated().sum()
```

```
[11]: 0
```

Check for and handle outliers.

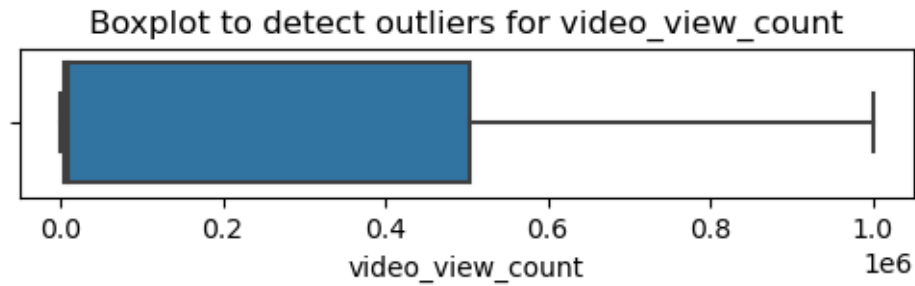
```
[12]: # Create a boxplot to visualize distribution of `video_duration_sec`
      ### YOUR CODE HERE ###
      plt.figure(figsize = (6,1))
      plt.title('Boxplot to detect outliers for video_duration_sec')

      sns.boxplot(x = data['video_duration_sec'])
      plt.show()
```



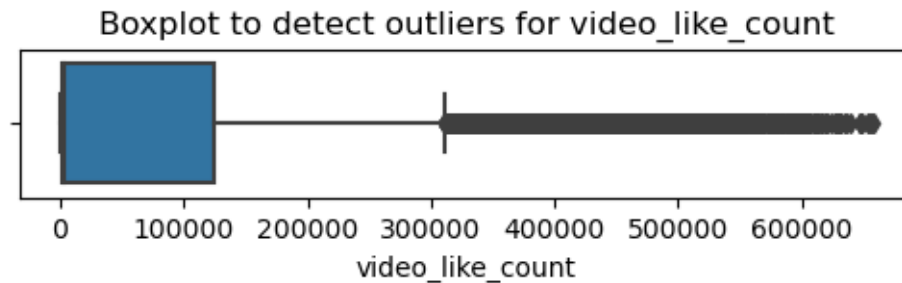
```
[13]: # Create a boxplot to visualize distribution of `video_view_count`
      ### YOUR CODE HERE ###
      plt.figure(figsize = (6,1))
      plt.title('Boxplot to detect outliers for video_view_count')

      sns.boxplot(x = data['video_view_count'])
      plt.show()
```



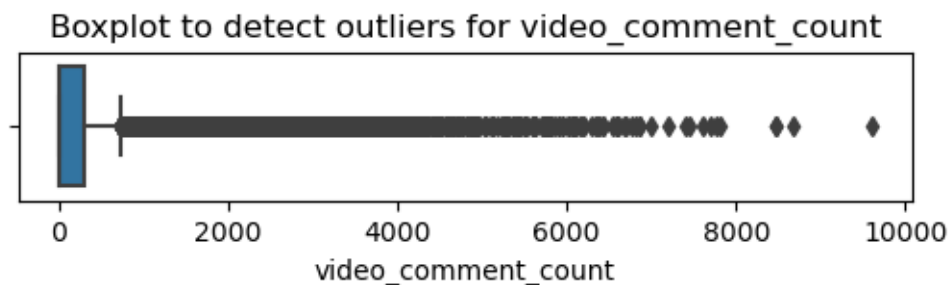
```
[14]: # Create a boxplot to visualize distribution of `video_like_count`
      ### YOUR CODE HERE ###
      plt.figure(figsize = (6,1))
      plt.title('Boxplot to detect outliers for video_like_count')

      sns.boxplot(x = data['video_like_count'])
      plt.show()
```



```
[15]: # Create a boxplot to visualize distribution of `video_comment_count`
      ### YOUR CODE HERE ###
      plt.figure(figsize = (6,1))
      plt.title('Boxplot to detect outliers for video_comment_count')

      sns.boxplot(x = data['video_comment_count'])
      plt.show()
```





```
[16]: # Check for and handle outliers for video_like_count
      ### YOUR CODE HERE ###

      percentile25 = data['video_like_count'].quantile(0.25)
      percentile75 = data['video_like_count'].quantile(0.75)

      iqr = percentile75 - percentile25

      upper_limit = percentile75 + (1.5 * iqr)

      data.loc[data['video_like_count'] > upper_limit, 'video_like_count'] =
        ↪upper_limit
```

Check class balance of the target variable. Remember, the goal is to predict whether the user of a given post is verified or unverified.

```
[17]: # Check class balance
      ### YOUR CODE HERE ###
      data['verified_status'].value_counts(normalize = True)
```

```
[17]: verified_status
      not verified    0.93712
      verified       0.06288
      Name: proportion, dtype: float64
```

Approximately 93.7% of the dataset represents videos posted by unverified accounts and 6.2% represents videos posted by verified accounts. So the outcome variable is not very balanced.

Use resampling to create class balance in the outcome variable, if needed.

```
[18]: # Use resampling to create class balance in the outcome variable, if needed

      # Identify data points from majority and minority classes
      ### YOUR CODE HERE ###
      data_majority = data[data["verified_status"] == "not verified"]
      data_minority = data[data["verified_status"] == "verified"]

      # Upsample the minority class (which is "verified")
      ### YOUR CODE HERE ###
      data_minority_upsampled = resample(data_minority,
                                         replace=True,
                                         n_samples=len(data_majority),
                                         random_state=0)

      # Combine majority class with upsampled minority class
      ### YOUR CODE HERE ###
```

```
data_upsampled = pd.concat([data_majority, data_minority_upsampled]).
    ↪reset_index(drop=True)

# Display new class counts
### YOUR CODE HERE ###
data_upsampled["verified_status"].value_counts()
```

```
[18]: verified_status
not verified    17884
verified       17884
Name: count, dtype: int64
```

Get the average video\_transcription\_text length for videos posted by verified accounts and the average video\_transcription\_text length for videos posted by unverified accounts.

```
[19]: # Get the average `video_transcription_text` length for claims and the average
    ↪`video_transcription_text` length for opinions
### YOUR CODE HERE ###
data_upsampled[["verified_status", "video_transcription_text"]].
    ↪groupby(by="verified_status")["video_transcription_text"].agg(func=lambda
    ↪array: np.mean([len(text) for text in array]))
```

```
[19]:          video_transcription_text
verified_status
not verified          89.401141
verified             84.569559
```

Extract the length of each video\_transcription\_text and add this as a column to the dataframe, so that it can be used as a potential feature in the model.

```
[20]: # Extract the length of each `video_transcription_text` and add this as a
    ↪column to the dataframe
### YOUR CODE HERE ###
data_upsampled["text_length"] = data_upsampled["video_transcription_text"].
    ↪apply(func=lambda text: len(text))
```

```
[21]: # Display first few rows of dataframe after adding new column
### YOUR CODE HERE ###
data_upsampled.head(5)
```

```
[21]:   # claim_status  video_id  video_duration_sec  \
0  1      claim  7017666017          59
1  2      claim  4014381136          32
2  3      claim  9859838091          31
3  4      claim  1866847991          25
4  5      claim  7105231098          19
```

	video_transcription_text	verified_status	\
0	someone shared with me that drone deliveries a...	not verified	
1	someone shared with me that there are more mic...	not verified	
2	someone shared with me that american industria...	not verified	
3	someone shared with me that the metro of st. p...	not verified	
4	someone shared with me that the number of busi...	not verified	

	author_ban_status	video_view_count	video_like_count	video_share_count	\
0	under review	343296.0	19425.0	241.0	
1	active	140877.0	77355.0	19034.0	
2	active	902185.0	97690.0	2858.0	
3	active	437506.0	239954.0	34812.0	
4	active	56167.0	34987.0	4110.0	

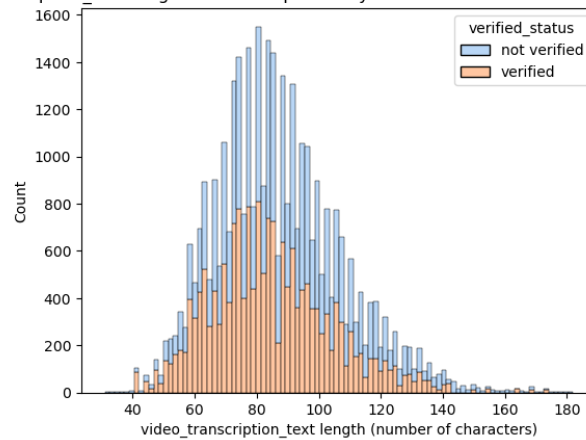
  

	video_download_count	video_comment_count	text_length
0	1.0	0.0	97
1	1161.0	684.0	107
2	833.0	329.0	137
3	1234.0	584.0	131
4	547.0	152.0	128

Visualize the distribution of `video_transcription_text` length for videos posted by verified accounts and videos posted by unverified accounts.

```
[22]: # Visualize the distribution of `video_transcription_text` length for videos
      ↪ posted by verified accounts and videos posted by unverified accounts
      # Create two histograms in one plot
      ### YOUR CODE HERE ###
      sns.histplot(data=data_upsampled, stat="count", multiple="stack",
      ↪ x="text_length", kde=False, palette="pastel",
      ↪ hue="verified_status", element="bars", legend=True)
      plt.title("Seaborn Stacked Histogram")
      plt.xlabel("video_transcription_text length (number of characters)")
      plt.ylabel("Count")
      plt.title("Distribution of video_transcription_text length for videos posted by
      ↪ verified accounts and videos posted by unverified accounts")
      plt.show()
```

Distribution of video\_transcription\_text length for videos posted by verified accounts and videos posted by unverified accounts



### 4.2.2 Task 2b. Examine correlations

Next, code a correlation matrix to help determine most correlated variables.

```
[23]: # Code a correlation matrix to help determine most correlated variables
      ### YOUR CODE HERE ###
      data_upsampled.corr(numeric_only = True)
```

```
[23]:
```

	#	video_id	video_duration_sec	\
#	1.000000	-0.000853	-0.011729	
video_id	-0.000853	1.000000	0.011859	
video_duration_sec	-0.011729	0.011859	1.000000	
video_view_count	-0.697007	0.002554	0.013589	
video_like_count	-0.626385	0.005993	0.004494	
video_share_count	-0.504015	0.010515	0.002206	
video_download_count	-0.487096	0.008753	0.003989	
video_comment_count	-0.413799	0.013983	-0.004586	
text_length	-0.193677	-0.007083	-0.002981	

	video_view_count	video_like_count	video_share_count	\
#	-0.697007	-0.626385	-0.504015	
video_id	0.002554	0.005993	0.010515	
video_duration_sec	0.013589	0.004494	0.002206	
video_view_count	1.000000	0.856937	0.711313	
video_like_count	0.856937	1.000000	0.832146	
video_share_count	0.711313	0.832146	1.000000	
video_download_count	0.690048	0.805543	0.710117	
video_comment_count	0.583485	0.686647	0.620182	
text_length	0.244693	0.216693	0.171651	

	video_download_count	video_comment_count	text_length
video_download_count	1.000000	0.690048	0.244693
video_comment_count	0.690048	1.000000	0.583485
text_length	0.244693	0.583485	1.000000

#	-0.487096	-0.413799	-0.193677
video_id	0.008753	0.013983	-0.007083
video_duration_sec	0.003989	-0.004586	-0.002981
video_view_count	0.690048	0.583485	0.244693
video_like_count	0.805543	0.686647	0.216693
video_share_count	0.710117	0.620182	0.171651
video_download_count	1.000000	0.857679	0.173396
video_comment_count	0.857679	1.000000	0.149750
text_length	0.173396	0.149750	1.000000

Visualize a correlation heatmap of the data.

```
[24]: # Create a heatmap to visualize how correlated variables are
      ### YOUR CODE HERE ###
      plt.figure(figsize=(8, 6))
      sns.heatmap(
          data_upsampled[["video_duration_sec", "claim_status", "author_ban_status",
          ↪ "video_view_count",
          ↪ "video_like_count", "video_share_count",
          ↪ "video_download_count", "video_comment_count", "text_length"]]
          .corr(numeric_only=True),
          annot=True,
          cmap="crest")
      plt.title("Heatmap of the dataset")
      plt.show()
```



One of the model assumptions for logistic regression is no severe multicollinearity among the features. Take this into consideration as you examine the heatmap and choose which features to proceed with.

**Question:** What variables are shown to be correlated in the heatmap?

**Answer:** The heatmap shows a strong correlation (0.86) between video\_view\_count and video\_like\_count. To ensure the logistic regression model satisfies the assumption of no severe multicollinearity, exclude video\_like\_count. Retain video\_view\_count, video\_share\_count, video\_download\_count, and video\_comment\_count as features.

### 4.3 PACE: Construct

After analysis and deriving variables with close relationships, it is time to begin constructing the model. Consider the questions in your PACE Strategy Document to reflect on the Construct stage.

#### 4.3.1 Task 3a. Select variables

Set your Y and X variables.

Select the outcome variable.

```
[25]: # Select outcome variable
      ### YOUR CODE HERE ###
      y = data_upsampled["verified_status"]
```

Select the features.

```
[26]: # Select features
      ### YOUR CODE HERE ###
      X = data_upsampled[["video_duration_sec", "claim_status", "author_ban_status",
      ↪ "video_view_count", "video_share_count", "video_download_count",
      ↪ "video_comment_count"]]

      # Display first few rows of features dataframe
      ### YOUR CODE HERE ###
      X.head()
```

```
[26]:  video_duration_sec  claim_status  author_ban_status  video_view_count  \
0           59          claim      under review          343296.0
1           32          claim           active          140877.0
2           31          claim           active          902185.0
3           25          claim           active          437506.0
4           19          claim           active          56167.0

      video_share_count  video_download_count  video_comment_count
0           241.0           1.0           0.0
1          19034.0          1161.0          684.0
2           2858.0           833.0          329.0
3          34812.0          1234.0          584.0
4           4110.0           547.0          152.0
```

#### 4.3.2 Task 3b. Train-test split

Split the data into training and testing sets.

```
[27]: # Split the data into training and testing sets
      ### YOUR CODE HERE ###
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
      ↪ random_state=0)
```

Confirm that the dimensions of the training and testing sets are in alignment.

```
[28]: # Get shape of each training and testing set
      ### YOUR CODE HERE ###
      X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[28]: ((26826, 7), (8942, 7), (26826,), (8942,))
```

### 4.3.3 Task 3c. Encode variables

Check the data types of the features.

```
[29]: # Check data types
      ### YOUR CODE HERE ###
      X_train.dtypes
```

```
[29]: video_duration_sec      int64
      claim_status           object
      author_ban_status      object
      video_view_count       float64
      video_share_count      float64
      video_download_count   float64
      video_comment_count    float64
      dtype: object
```

```
[30]: # Get unique values in `claim_status`
      ### YOUR CODE HERE ###
      X_train["claim_status"].unique()
```

```
[30]: array(['opinion', 'claim'], dtype=object)
```

```
[31]: # Get unique values in `author_ban_status`
      ### YOUR CODE HERE ###
      X_train["author_ban_status"].unique()
```

```
[31]: array(['active', 'under review', 'banned'], dtype=object)
```

As shown above, the `claim_status` and `author_ban_status` features are each of data type `object` currently. In order to work with the implementations of models through `sklearn`, these categorical features will need to be made numeric. One way to do this is through one-hot encoding.

Encode categorical features in the training set using an appropriate method.

```
[32]: # Select the training features that needs to be encoded
      ### YOUR CODE HERE ###
      X_train_to_encode = X_train[["claim_status", "author_ban_status"]]

      # Display first few rows
      ### YOUR CODE HERE ###
      X_train_to_encode.head()
```

```
[32]:      claim_status  author_ban_status
33058      opinion      active
20491      opinion      active
25583      opinion      active
18474      opinion      active
27312      opinion      active
```



```
[33]: # Set up an encoder for one-hot encoding the categorical features
      ### YOUR CODE HERE ###
      X_encoder = OneHotEncoder(drop='first', sparse_output=False)
```

```
[34]: # Fit and transform the training features using the encoder
      ### YOUR CODE HERE ###
      X_train_encoded = X_encoder.fit_transform(X_train_to_encode)
```

```
[35]: # Get feature names from encoder
      ### YOUR CODE HERE ###
      X_encoder.get_feature_names_out()
```

```
[35]: array(['claim_status_opinion', 'author_ban_status_banned',
        'author_ban_status_under review'], dtype=object)
```

```
[36]: # Display first few rows of encoded training features
      ### YOUR CODE HERE ###
      X_train_encoded
```

```
[36]: array([[1., 0., 0.],
        [1., 0., 0.],
        [1., 0., 0.],
        ...,
        [1., 0., 0.],
        [1., 0., 0.],
        [0., 1., 0.]])
```

```
[37]: # Place encoded training features (which is currently an array) into a dataframe
      ### YOUR CODE HERE ###
      X_train_encoded_df = pd.DataFrame(data=X_train_encoded, columns=X_encoder.
      ↪get_feature_names_out())

      # Display first few rows
      ### YOUR CODE HERE ###
      X_train_encoded_df.head()
```

```
[37]:   claim_status_opinion  author_ban_status_banned \
0                1.0                0.0
1                1.0                0.0
2                1.0                0.0
3                1.0                0.0
4                1.0                0.0

      author_ban_status_under review
0                0.0
1                0.0
2                0.0
```

```
3          0.0
4          0.0
```

```
[38]: # Display first few rows of `X_train` with `claim_status` and
      ↪ `author_ban_status` columns dropped (since these features are being
      ↪ transformed to numeric)
      ### YOUR CODE HERE ###
      X_train.drop(columns=["claim_status", "author_ban_status"]).head()
```

```
[38]:      video_duration_sec  video_view_count  video_share_count  \
33058              33          2252.0             23.0
20491              52          6664.0            550.0
25583              37          6327.0            257.0
18474              57          1702.0             28.0
27312              21          3842.0            101.0

      video_download_count  video_comment_count
33058                  4.0                  0.0
20491                 53.0                  2.0
25583                  3.0                  0.0
18474                  0.0                  0.0
27312                  1.0                  0.0
```

```
[39]: # Concatenate `X_train` and `X_train_encoded_df` to form the final dataframe
      ↪ for training data (`X_train_final`)
      # Note: Using `.reset_index(drop=True)` to reset the index in `X_train` after
      ↪ dropping `claim_status` and `author_ban_status`,
      # so that the indices align with those in `X_train_encoded_df` and `count_df`
      ### YOUR CODE HERE ###
      X_train_final = pd.concat([X_train.drop(columns=["claim_status",
      ↪ "author_ban_status"]), X_train_encoded_df], axis=1)

      # Display first few rows
      ### YOUR CODE HERE ###
      X_train_final = pd.concat([X_train.drop(columns=["claim_status",
      ↪ "author_ban_status"]), X_train_encoded_df], axis=1)
```

Check the data type of the outcome variable.

```
[40]: # Check data type of outcome variable
      ### YOUR CODE HERE ###
      y_train.dtype
```

```
[40]: dtype('O')
```

```
[41]: # Get unique values of outcome variable
      ### YOUR CODE HERE ###
```

```
y_train.unique()
```

```
[41]: array(['verified', 'not verified'], dtype=object)
```

As shown above, the outcome variable is of data type `object` currently. One-hot encoding can be used to make this variable numeric.

Encode categorical values of the outcome variable the training set using an appropriate method.

```
[42]: # Set up an encoder for one-hot encoding the categorical outcome variable
      ### YOUR CODE HERE ###
      y_encoder = OneHotEncoder(drop='first', sparse_output=False)
```

```
[43]: # Encode the training outcome variable
      # Notes:
      #   - Adjusting the shape of `y_train` before passing into `.fit_transform()`,
      #     ↳ since it takes in 2D array
      #   - Using `.ravel()` to flatten the array returned by `.fit_transform()`, so
      #     ↳ that it can be used later to train the model
      ### YOUR CODE HERE ###
      y_train_final = y_encoder.fit_transform(y_train.values.reshape(-1, 1)).ravel()

      # Display the encoded training outcome variable
      ### YOUR CODE HERE ###
      y_train_final
```

```
[43]: array([1., 1., 1., ..., 1., 1., 0.])
```

#### 4.3.4 Task 3d. Model building

Construct a model and fit it to the training set.

```
[44]: # Construct a logistic regression model and fit it to the training set
      ### YOUR CODE HERE ###
      log_clf = LogisticRegression(random_state=0, max_iter=800).fit(X_train_final,
      ↳ y_train_final)
```

#### 4.4 PACE: Execute

Consider the questions in your PACE Strategy Document to reflect on the Execute stage.

##### 4.4.1 Task 4a. Results and evaluation

Evaluate your model.

Encode categorical features in the testing set using an appropriate method.

```
[45]: # Select the testing features that needs to be encoded
      ### YOUR CODE HERE ###
```

```
X_test_to_encode = X_test[["claim_status", "author_ban_status"]]

# Display first few rows
### YOUR CODE HERE ###
X_test_to_encode.head()
```

```
[45]:      claim_status  author_ban_status
21061      opinion      active
31748      opinion      active
20197      claim      active
5727       claim      active
11607      opinion      active
```

```
[46]: # Transform the testing features using the encoder
### YOUR CODE HERE ###
X_test_encoded = X_encoder.transform(X_test_to_encode)

# Display first few rows of encoded testing features
### YOUR CODE HERE ###
X_test_encoded
```

```
[46]: array([[1., 0., 0.],
        [1., 0., 0.],
        [0., 0., 0.],
        ...,
        [1., 0., 0.],
        [0., 0., 1.],
        [1., 0., 0.]])
```

```
[47]: # Place encoded testing features (which is currently an array) into a dataframe
### YOUR CODE HERE ###
X_test_encoded_df = pd.DataFrame(data=X_test_encoded, columns=X_encoder.
    ↪get_feature_names_out())

# Display first few rows
### YOUR CODE HERE ###
X_test_encoded_df.head()
```

```
[47]:      claim_status_opinion  author_ban_status_banned  \
0                1.0                0.0
1                1.0                0.0
2                0.0                0.0
3                0.0                0.0
4                1.0                0.0

      author_ban_status_under review
0                0.0
```

```

1          0.0
2          0.0
3          0.0
4          0.0

```

```

[48]: # Display first few rows of `X_test` with `claim_status` and
      ↪ `author_ban_status` columns dropped (since these features are being
      ↪ transformed to numeric)
      ### YOUR CODE HERE ###
      X_test.drop(columns=["claim_status", "author_ban_status"]).head()

```

```

[48]:      video_duration_sec  video_view_count  video_share_count  \
21061          41          2118.0          57.0
31748          27          5701.0          157.0
20197          31         449767.0        75385.0
5727           19        792813.0       56597.0
11607          54          2044.0          68.0

      video_download_count  video_comment_count
21061           5.0           2.0
31748           1.0           0.0
20197        5956.0        1789.0
5727        5146.0        3413.0
11607        19.0           2.0

```

```

[49]: # Concatenate `X_test` and `X_test_encoded_df` to form the final dataframe for
      ↪ training data (`X_test_final`)
      # Note: Using `.reset_index(drop=True)` to reset the index in `X_test` after
      ↪ dropping `claim_status`, and `author_ban_status`,
      # so that the indices align with those in `X_test_encoded_df` and
      ↪ `test_count_df`
      ### YOUR CODE HERE ###
      X_test_final = pd.concat([X_test.drop(columns=["claim_status",
      ↪ "author_ban_status"]), X_test_encoded_df], axis=1)

      # Display first few rows
      ### YOUR CODE HERE ###
      X_test_final.head()

```

```

[49]:      video_duration_sec  video_view_count  video_share_count  \
0          41          2118.0          57.0
1          27          5701.0          157.0
2          31         449767.0        75385.0
3          19        792813.0       56597.0
4          54          2044.0          68.0

      video_download_count  video_comment_count  claim_status_opinion  \

```

0	5.0	2.0	1.0
1	1.0	0.0	1.0
2	5956.0	1789.0	0.0
3	5146.0	3413.0	0.0
4	19.0	2.0	1.0

	author_ban_status_banned	author_ban_status_under review
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

Test the logistic regression model. Use the model to make predictions on the encoded testing set.

```
[50]: # Use the logistic regression model to get predictions on the encoded testing set
      ↪set
      ### YOUR CODE HERE ###
      y_pred = log_clf.predict(X_test_final)
```

Display the predictions on the encoded testing set.

```
[51]: # Display the predictions on the encoded testing set
      ### YOUR CODE HERE ###
      print(y_pred)
```

```
[1.  1.  0. ... 1.  0.  1.]
```

Display the true labels of the testing set.

```
[52]: # Display the true labels of the testing set
      ### YOUR CODE HERE ###
      print(y_test)
```

```
21061    verified
31748    verified
20197    verified
5727     not verified
11607     not verified
...
14756     not verified
26564     verified
14800     not verified
35705     verified
31060     verified
```

```
Name: verified_status, Length: 8942, dtype: object
```

Encode the true labels of the testing set so it can be compared to the predictions.

```
[53]: # Encode the testing outcome variable
# Notes:
# - Adjusting the shape of `y_test` before passing into `.transform()`, since
  ↳ it takes in 2D array
# - Using `.ravel()` to flatten the array returned by `.transform()`, so that
  ↳ it can be used later to compare with predictions
### YOUR CODE HERE ###
y_test_final = y_encoder.transform(y_test.values.reshape(-1, 1)).ravel()

# Display the encoded testing outcome variable
print(y_test_final)
```

```
[1. 1. 1. ... 0. 1. 1.]
```

Confirm again that the dimensions of the training and testing sets are in alignment since additional features were added.

```
[54]: # Get shape of each training and testing set
### YOUR CODE HERE ###
X_train_final.shape, y_train_final.shape, X_test_final.shape, y_test_final.shape
```

```
[54]: ((26826, 8), (26826,), (8942, 8), (8942,))
```

#### 4.4.2 Task 4b. Visualize model results

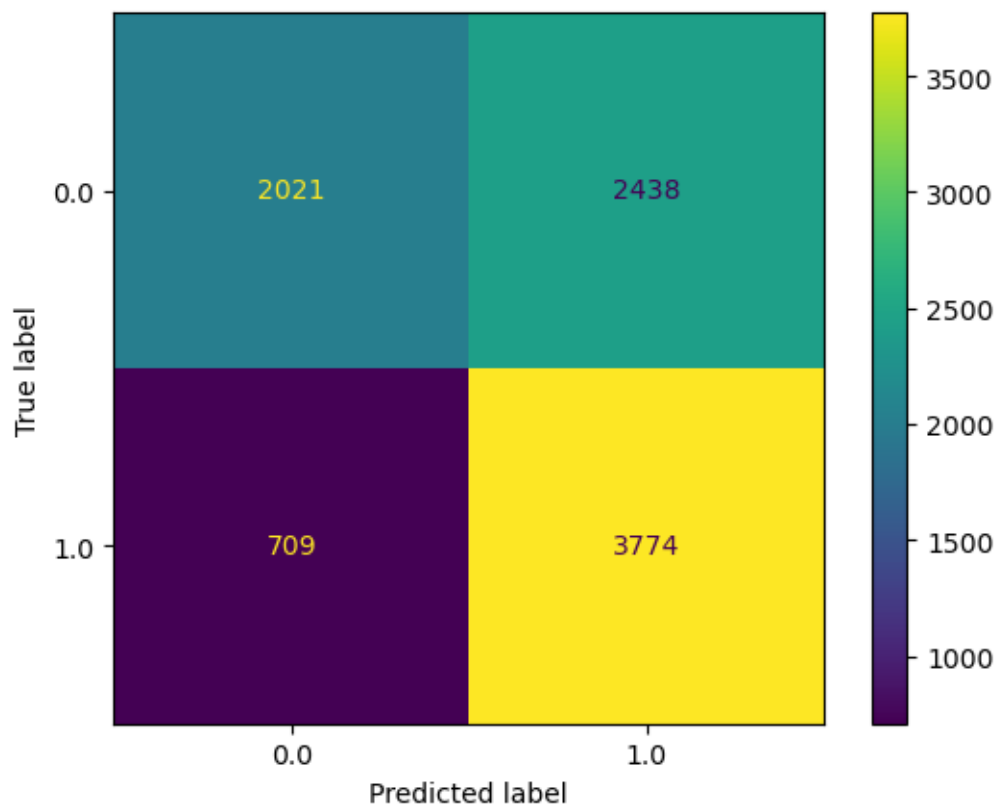
Create a confusion matrix to visualize the results of the logistic regression model.

```
[55]: # Compute values for confusion matrix
### YOUR CODE HERE ###
log_cm = confusion_matrix(y_test_final, y_pred, labels=log_clf.classes_)

# Create display of confusion matrix
### YOUR CODE HERE ###
log_disp = ConfusionMatrixDisplay(confusion_matrix=log_cm,
  ↳ display_labels=log_clf.classes_)

# Plot confusion matrix
### YOUR CODE HERE ###
log_disp.plot()

# Display plot
### YOUR CODE HERE ###
plt.show()
```



Create a classification report that includes precision, recall, f1-score, and accuracy metrics to evaluate the performance of the logistic regression model.

```
[56]: # Create a classification report
      ### YOUR CODE HERE ###
      target_labels = ["verified", "not verified"]
      print(classification_report(y_test_final, y_pred, target_names=target_labels))
```

	precision	recall	f1-score	support
verified	0.74	0.45	0.56	4459
not verified	0.61	0.84	0.71	4483
accuracy			0.65	8942
macro avg	0.67	0.65	0.63	8942
weighted avg	0.67	0.65	0.63	8942



#### 4.4.3 Task 4c. Interpret model coefficients

```
[57]: # Get the feature names from the model and the model coefficients (which
      ↪ represent log-odds ratios)
      # Place into a DataFrame for readability
      ### YOUR CODE HERE ###
      pd.DataFrame(data={"Feature Name":log_clf.feature_names_in_, "Model_
      ↪Coefficient":log_clf.coef_[0]})
```

```
[57]:
```

	Feature Name	Model Coefficient
0	video_duration_sec	8.493546e-03
1	video_view_count	-2.277453e-06
2	video_share_count	5.458611e-06
3	video_download_count	-2.143023e-04
4	video_comment_count	3.899371e-04
5	claim_status_opinion	3.772015e-04
6	author_ban_status_banned	-1.675961e-05
7	author_ban_status_under review	-7.084767e-07

#### 4.4.4 Task 4d. Conclusion

1. What are the key takeaways from this project?
2. What results can be presented from this project?

**Answer:**

1. Key takeaways:
  - Dropped video\_like\_count to prevent multicollinearity issues.
  - Additional second of video raises log-odds of user verification by 0.009.
  - Precision: 61%, recall: 84%, overall accuracy: barely acceptable.
2. Our logistic regression model for predicting verified status based on video features showed decent predictive power. Longer videos are more likely to result in user verification, whereas other video features have small coefficients and less significant associations.

**Congratulations!** You've completed this lab. However, you may not notice a green check mark next to this item on Coursera's platform. Please continue your progress regardless of the check mark. Just click on the "save" icon at the top of this notebook to ensure your work has been logged.